


W

UNIVERSITY of WA

# What goes around comes around




- Pre-2004: commercial RDBMS, some open source
- 2004 Dean et al. MapReduce
- 2008 Hadoop 0.17 release
- 2008 Olston et al. Pig: Relational Algebra on Hadoop
- 2008 DryadLINQ: Relational Algebra in a Hadoop-like system
- 2009 Thusoo et al. HIVE: SQL on Hadoop
- 2009 Hbase: Indexing for Hadoop
- 2010 Dietrich et al. Schemas and Indexing for Hadoop
- 2012 Transactions in HBase (plus VoltDB, other NewSQL systems)
- But also some permanent contributions:
  - Fault tolerance
  - Schema-on-Read
  - User-defined functions that don't suck

4/9/2013

Bill Howe, UW

2



02:17 / 14:02



Quick Question

Home

Video Lectures

Discussion For

Programming A

Quizzes

Peer Assessme

Optional Real-v

Help Center

